	Course	Databases and Information Systems 2020		
	Exercise Sheet	3		
	Points	–		
	Release Date	May 12th 2020	Due Date	May 19th 2020

3 Synchronization with Locking Protocols

Note

- Include SQL commands in your report and not as separate SQL files.

3.1 Isolation Levels and SQL

- a) Open a connection to database **dis-2020**¹ with a tool of your choice. If you prefer, you can use a local PostgreSQL installation.
- How can you determine the currently set isolation level?
 - What is the default isolation level of PostgreSQL?
 - How can the isolation level be changed during a session in PostgreSQL?

Provide a solution applicable to all kinds of tools. Avoid tool specific shortcuts or QoL features.
Hint: Consider checking the PostgreSQL documentation.

- b) Write down a SQL code for creating a simple table **OPK** with the columns **ID** and **NAME**. No column is declared as a primary key.
- c) Write down a SQL command suitable for filling **OPK** with data. Afterward, the data should look like this:

OPK

ID	NAME
1	shaggy
2	fred
3	velma
4	scooby
5	daphne

- d) Consider an open connection to an arbitrary database like **DB2**² with a current transaction level of **Read Committed** (PostgreSQL uses a Multiversion Concurrency Control mechanism, which makes this question not applicable). The connection executes the following SQL commands:

```
-- Begin Transaction
START TRANSACTION;
-- Query one row from table OPK
SELECT * FROM OPK WHERE ID = 3;
-- ! Check for current locks !
COMMIT;
```


Discuss what locks you would expect are held at this point (and before a commit)?

Hint: You can use Slides 8-11 from the Additional Slides for this Exercise Sheet. Caution: The Slides were originally written for the database DB2.

- e) Discuss what locks you would expect are held if the SQL command from d) is done with the isolation level **REPEATABLE READ**.

¹Database from sheet 1 & 2

²Consider looking into the Additional Slides (Slides 8-11).

	Course	Databases and Information Systems 2020		
	Exercise Sheet	3		
	Points	–		
	Release Date	May 12th 2020	Due Date	May 19th 2020

3.2 Lock Conflicts

- a) Given the following scenario: You have two open connections. Connection 1 (isolation level *Repeatable Read*) queries the rows with *ID>3* from *OPK* but does not commit yet. Connection 2 (isolation level *Committed Read*) adds a new row to the table that satisfies the selection predicate.

- What happens? What is the output of Connection 1?
- Compare the state before the transactions with the state after the transactions.
- What can be observed if Connection 1 commits and execute its SQL command again?
- Can we observe a Canonical Synchronization Problem? If yes, explain which one and why it appears.


T ₁ (Repeatable Read)	T ₂ (Committed Read)
<pre>START TRANSACTION; SELECT * FROM OPK WHERE ID > 3; COMMIT;</pre>	<pre>START TRANSACTION; INSERT INTO OPK VALUES(6, 'scrappy'); COMMIT;</pre>

- b) Repeat the scenario from a) with the isolation level of Connection 1 set to *Serializable*. Now repeat the steps from a).

- What can be observed now? What is the output of Connection 1?
- Compare the state before the transactions with the state after the transactions.
- Can we observe a Canonical Synchronization Problems? If yes, explain which one and why it appears.

- c) Given the following scenario: You have two open connections. Connection 1 (isolation level *Serializable*) queries one row from table *OPK* (selection via the *ID* attribute, *ID=3*) but does not commit yet. Connection 2 (isolation level *Committed Read*) change another row's *NAME* value.

- In this scenario Connection 2 has to wait until Connection 1 commits. Explain why.
- Discuss, what lock can be potentially encountered on the table *OPK*? Which Connection do the locks belong to?

	Course	Databases and Information Systems 2020		
	Exercise Sheet	3		
	Points	–		
	Release Date	May 12th 2020	Due Date	May 19th 2020

T ₁ (Serializable)	T ₂ (Committed Read)
<pre> START TRANSACTION; SELECT * FROM OPK WHERE ID = 3; COMMIT; </pre>	<pre> START TRANSACTION; UPDATE OPK SET NAME = 'norville' WHERE id = 1; COMMIT; </pre>

- d) Write down a SQL code for creating a second table MPK, again with columns ID and Name, and declare ID a primary key. The new table should be filled with the same initial data as the first table OPK. The SQL code for filling MPK has to be added, too.
- e) The transactions from c) are repeated with the table MPK. This time, Connection 2 doesn't have to wait for the commit of Connection 1. Explain the behavior.