

Exercise 3 Report

Gruppe 16

Anastasiia Rubanovych

Sebastian Funck

May 17, 2020

3.1 Isolation Levels and SQL

a)

- **How can you determine the currently set isolation level?**

The default isolation level of a database is configured inside the `postgresql.conf` with the attribute `default_transaction_isolation`. It can be queried via `SHOW default_transaction_isolation`.

- **What is the default isolation level of PostgreSQL?**

The default isolation level of PostgreSQL is `read committed`.

- **How can the isolation level be changed during a session in PostgreSQL?**

```
SET SESSION CHARACTERISTICS AS TRANSACTION { SERIALIZABLE | REPEATABLE  
READ | READ COMMITTED | READ UNCOMMITTED }
```

b)

```
create table OPK (  
  ID    int4 ,  
  NAME  varchar(64)  
)
```

c)

```
insert into OPK (ID, NAME)      VALUES (1, 'shaggy'), (2, 'fred'),      (3, 'shaggy')
```

d)

- **Discuss what locks you would expect are held at this point (and before a commit) with Read Committed?**

We would expect 1 table lock and 1 row lock to be hold before and 0 table locks and 0 row locks to be hold after the commit. Cursor Stability (Read Committed) never holds more than 1 row lock.

- **Discuss what locks you would expect are held at this point (and before a commit) with Repeatable Read?**

We would expect 1 table lock to be hold before and 0 locks after the commit. Repeatable Read locks the complete table until commit.

3.2 Lock Conflicts

a)

- **What happens? What is the output of Connection 1?**

The output of Connection 1 is

```
(4, 'scooby ')  
(5, 'daphne ')
```

- **Compare the state before the transactions with the state after the transactions.**

After the transactions, one new row has been added to the OPK table.

- **What can be observed if Connection 1 commits and execute its SQL command again?**

The output of Connection 1 then becomes

```
(4, 'scooby ')  
(5, 'daphne ')  
(6, 'scrappy ')
```

- **Can we observe a Canonical Synchronization Problem? If yes, explain which one and why it appears.**

We cannot observe any canonical synchronization problems.

b)

- **What happens? What is the output of Connection 1?**

The output of Connection 1 is

```
(4, 'scooby ')  
(5, 'daphne ')
```

- **Compare the state before the transactions with the state after the transactions.**

After the transactions, one new row has been added to the OPK table.

- **What can be observed if Connection 1 commits and execute its SQL command again?**

The output of Connection 1 then becomes

```
(4, 'scooby')
(5, 'daphne')
(6, 'scrappy')
```

- **Can we observe a Canonical Synchronization Problem? If yes, explain which one and why it appears.**

We cannot observe any canonical synchronization problems.

c)

- **In this scenario Connection 2 has to wait until Connection 1 commits. Explain why.**

The **Serializable** isolation level ensures a serial execution of transactions, therefore Connection 1 has to be finished before Connection 2 can continue.

- **Discuss, what lock can be potentially encountered on the table OPK? Which Connection do the locks belong to?**

Connection 1 is likely to hold an R lock on the table.

d)

```
create table MPK (
    id int4,
    name varchar(64),
    CONSTRAINT mpk_pk PRIMARY KEY (id)
);

insert into MPK (ID, NAME)
VALUES (1, 'shaggy'), (2, 'fred'),
(3, 'velma'), (4, 'scooby'), (5, 'daphne');
```

e)

We observed the exact same behaviour as for *c*).