

Exercise 3 Report

Gruppe 16

Anastasiia Rubanovych

Sebastian Funck

May 19, 2020

3.1 Isolation Levels and SQL

a)

- **How can you determine the currently set isolation level?**

The default isolation level of a database is configured inside the `postgresql.conf` with the attribute `default_transaction_isolation`. It can be queried via `SHOW default_transaction_isolation`.

- **What is the default isolation level of PostgreSQL?**

The default isolation level of PostgreSQL is `read committed`.

- **How can the isolation level be changed during a session in PostgreSQL?**

```
SET SESSION CHARACTERISTICS AS TRANSACTION { SERIALIZABLE | REPEATABLE  
READ | READ COMMITTED | READ UNCOMMITTED }
```

b)

```
create table OPK (  
  ID    int4 ,  
  NAME  varchar(64)  
)
```

c)

```
insert into OPK (ID, NAME)  
VALUES (1, 'shaggy'), (2, 'fred'),  
       (3, 'velma'), (4, 'scooby'), (5, 'daphne');
```

d)

- **Discuss what locks you would expect are held at this point (and before a commit) with Read Committed?**

We would expect 1 shared table lock and 1 row lock to be hold before and 0 table locks and 0 row locks to be hold after the commit. Cursor Stability (Read Committed) never holds more than 1 row lock.

- **Discuss what locks you would expect are held at this point (and before a commit) with Repeatable Read?**

We would expect 1 table lock to be hold before and 0 locks after the commit. Repeatable Read locks the complete table until commit.

3.2 Lock Conflicts

a)

- **What happens? What is the output of Connection 1?**

While Connection 1 is executing the query, Connection 2 inserts a row that matches the query of Connection 1. This results in Connection 1 committing after Connection 2, but will not select the newly inserted item. Therefore the output of Connection 1 is

```
(4, 'scooby')
(5, 'daphne')
```

- **Compare the state before the transactions with the state after the transactions.**

After the commit of Connection 2 one new row has been added to the OPK table. This happens while Connection 1 is still executing.

- **What can be observed if Connection 1 commits and execute its SQL command again?**

The output of Connection 1 becomes

```
(4, 'scooby')
(5, 'daphne')
(6, 'scrappy')
```

- **Can we observe a Canonical Synchronization Problem? If yes, explain which one and why it appears.**

Yes, we can observe a Phantom Read because Connection 1 is querying a range while Connection 2 inserts a record (6, `scrappy`) into the table which would match that queried range.

b)

- **What happens? What is the output of Connection 1?**

Because Connection 1 is set to isolation mode **Serializable** the scheduler has to ensure that this transactions is executed in a serializable linear order. Therefore Connection 2 has to wait for Transaction 1 to finish.

```
(4, 'scooby')
(5, 'daphne')
```

- **Compare the state before the transactions with the state after the transactions.**

After the transaction of Connection 1 there are only 5 rows in the data set. Only after connection 2 has been completed, the 6th row is added.

- **What can be observed if Connection 1 commits and execute its SQL command again?**

The output of Connection 1 stays the same

```
(4, 'scooby')
(5, 'daphne')
(6, 'scrappy')
```

- **Can we observe a Canonical Synchronization Problem? If yes, explain which one and why it appears.**

By definition of isolation level **Serializable** there can't occur any canonical synchronization problems.

c)

- **In this scenario Connection 2 has to wait until Connection 1 commits. Explain why.**

ID=1 is in the context of table OKP a range based query, because multiple rows could match this criteria. In addition Connection 1 is in isolation level **Serializable**. This results in Connection 1 holding an R-lock on the complete table OKP and Connection 2 waiting for Connection 1 to commit.

- **Discuss, what lock can be potentially encountered on the table OPK? Which Connection do the locks belong to?**

Connection 1 is likely to hold an R lock on the table OKP.

d)

```
create table MPK (
    id int4,
    name varchar(64),
```

```

        CONSTRAINT mpk_pk PRIMARY KEY (id)
    );

insert into MPK (ID, NAME)
    VALUES (1, 'shaggy'), (2, 'fred'),
        (3, 'velma'), (4, 'scooby'), (5, 'daphne');

```

e)

With ID being a **PRIMARY KEY**, every row is uniquely identifiable by its id. This means the query of Connection 1 is no longer range-query and enables Connection 1 to hold a narrower lock (R-lock on row ID:3). On the otherhand Connection 2 now can hold an R-lock on row ID:3 while Connection 1 is executing and doesn't have to wait for it to finish.